

# Exercise 3-2

Author: Barbara Gaskins

Date: 11-02-2024

## Exercise 1 - 1

```
In [1]: from os.path import basename, exists
```

```
def download(url):
    filename = basename(url)
    if not exists(filename):
        from urllib.request import urlretrieve

        local, _ = urlretrieve(url, filename)
        print("Downloaded " + local)
```

```
download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/thinl
download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/thinl
```

```
In [2]: download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/nsfg
```

```
download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/2002l
download(
    "https://github.com/AllenDowney/ThinkStats2/raw/master/code/2002FemPr
)
```

```
In [3]: import nsfg
```

```
In [4]: preg = nsfg.ReadFemPreg()
preg.head()
```

```
Out[4]:
```

	caseid	pregordr	howpreg_n	howpreg_p	moscurrp	nowprgdk	pregend1	pregend2
0	1	1	NaN	NaN	NaN	NaN	6.0	NaN
1	1	2	NaN	NaN	NaN	NaN	6.0	NaN
2	2	1	NaN	NaN	NaN	NaN	5.0	NaN
3	2	2	NaN	NaN	NaN	NaN	6.0	NaN
4	2	3	NaN	NaN	NaN	NaN	6.0	NaN

5 rows × 244 columns

Select the `birthord` column, print the value counts, and compare to results published in the [c](#)

```
In [5]: preg.birthord.value_counts().sort_index()
```

```
Out[5]: 1.0    4413
        2.0    2874
        3.0    1234
        4.0     421
        5.0    126
        6.0     50
        7.0     20
        8.0      7
        9.0      2
       10.0      1
        Name: birthord, dtype: int64
```

We can also use `isnull` to count the number of nans.

```
In [6]: preg.birthord.isnull().sum()
```

```
Out[6]: 4445
```

Select the `prglngth` column, print the value counts, and compare to results published in the [c](#)

```
In [7]: preg.prglngth.value_counts().sort_index()
```

```
Out[7]: 0      15
        1       9
        2      78
        3     151
        4     412
        5     181
        6     543
        7     175
        8     409
        9     594
       10     137
       11     202
       12     170
       13     446
       14      29
       15      39
       16      44
       17     253
       18      17
       19      34
       20      18
       21      37
       22     147
       23      12
       24      31
       25      15
       26     117
       27       8
       28      38
       29      23
       30     198
       31      29
       32     122
       33      50
       34      60
       35     357
       36     329
       37     457
       38     609
       39    4744
       40    1120
       41     591
       42     328
       43     148
       44      46
       45      10
       46       1
       47       1
       48       7
       50       2
      Name: prglngth, dtype: int64
```

To compute the mean of a column, you can invoke the `mean` method on a Series. For example, the mean birthweight in pounds:

```
In [8]: preg.totalwgt_lb.mean()
```

```
Out[8]: 7.265628457623368
```

Create a new column named `totalwgt_kg` that contains birth weight in kilograms. Compute its mean. Remember that when you create a new column, you have to use dictionary syntax, not dot notation.

```
In [9]: preg['totalwgt_kg'] = preg.totalwgt_lb / 2.2
preg.totalwgt_kg.mean()
```

```
Out[9]: 3.302558389828807
```

`nsfg.py` also provides `ReadFemResp`, which reads the female respondents file and returns a `DataFrame`:

```
In [10]: download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/2002I")
download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/2002I")
```

```
In [11]: resp = nsfg.ReadFemResp()
```

`DataFrame` provides a method `head` that displays the first five rows:

```
In [12]: resp.head()
```

```
Out[12]:
```

	caseid	rscrinf	rdormres	rostscrn	rscreenhisp	rscreenrace	age_a	age_r	cmbirth
0	2298	1	5	5	1	5.0	27	27	902
1	5012	1	5	1	5	5.0	42	42	718
2	11586	1	5	1	5	5.0	43	43	708
3	6794	5	5	4	1	5.0	15	15	1042
4	616	1	5	4	1	5.0	20	20	991

5 rows × 3087 columns

Select the `age_r` column from `resp` and print the value counts. How old are the youngest and oldest respondents?

```
In [13]: resp.age_r.value_counts().sort_index()
```

```
Out[13]: 15    217
         16    223
         17    234
         18    235
         19    241
         20    258
         21    267
         22    287
         23    282
         24    269
         25    267
         26    260
         27    255
         28    252
         29    262
         30    292
         31    278
         32    273
         33    257
         34    255
         35    262
         36    266
         37    271
         38    256
         39    215
         40    256
         41    250
         42    215
         43    253
         44    235
Name: age_r, dtype: int64
```

We can use the `caseid` to match up rows from `resp` and `preg`. For example, we can select row from `resp` for `caseid` 2298 like this:

```
In [14]: resp[resp.caseid == 2298]
```

```
Out[14]:   caseid  rscrinf  rdormres  rostscrn  rscreenhisp  rscreenrace  age_a  age_r  cmbirth
0    2298         1         5         5             1           5.0      27      27      902
```

1 rows × 3087 columns

And we can get the corresponding rows from `preg` like this:

```
In [15]: preg[preg.caseid == 2298]
```

```
Out[15]:   caseid  pregordr  howpreg_n  howpreg_p  moscurrp  nowprgdk  pregend1  pregenr
2610    2298         1         NaN         NaN         NaN         NaN         6.0      NaN
2611    2298         2         NaN         NaN         NaN         NaN         6.0      NaN
2612    2298         3         NaN         NaN         NaN         NaN         6.0      NaN
2613    2298         4         NaN         NaN         NaN         NaN         6.0      NaN
```

4 rows × 245 columns

How old is the respondent with `caseid` 1?

```
In [16]: resp[resp.caseid == 1].age_r
```

```
Out[16]: 1069    44
         Name: age_r, dtype: int64
```

What are the pregnancy lengths for the respondent with `caseid` 2298?

```
In [17]: preg[preg.caseid == 2298].prglngth
```

```
Out[17]: 2610    40
         2611    36
         2612    30
         2613    40
         Name: prglngth, dtype: int64
```

What was the birthweight of the first baby born to the respondent with `caseid` 5012?

```
In [18]: preg[preg.caseid == 5012].birthwgt_lb
```

```
Out[18]: 5515    6.0
         Name: birthwgt_lb, dtype: float64
```

## Exercise 1 - 2

Print the value counts for the variable `pregnum`

```
In [19]: resp_pregnum_counts = resp.pregnum.value_counts().sort_index()
         resp_pregnum_counts
```

```
Out[19]: 0      2610
         1      1267
         2      1432
         3      1110
         4       611
         5       305
         6       150
         7        80
         8        40
         9        21
        10         9
        11         3
        12         2
        14         2
        19         1
         Name: pregnum, dtype: int64
```

```
In [20]: # Get the number of caseids using value count and convert it to a list
         # Using the Collections package get the frequency of the number of births
```

```
import collections
preg_count = preg['caseid'].value_counts().tolist()

preg_pregnum = collections.Counter(preg_count)
preg_pregnum
```

```
Out[20]: Counter({19: 1,
                  14: 2,
                  12: 2,
                  11: 3,
                  10: 9,
                  9: 21,
                  8: 40,
                  7: 80,
                  6: 150,
                  5: 305,
                  4: 611,
                  3: 1110,
                  2: 1432,
                  1: 1267})
```

## Exercise 2 - 1

For an evening news segment, I would focus on presenting the average and range of pregnancy. The average is significant because most people associate pregnancy with the standard 40-week. Highlighting the range is equally important, as it reveals that not all babies are born exactly at 40. This information addresses the curiosity of viewers who wonder about the earliest and latest weeks babies can be born.

For an expectant parent feeling anxious, I would emphasize the average as a reassuring reference, offering a sense of predictability during a time of uncertainty.

From a factual standpoint, there is a common belief that first babies are more likely to be "late," while subsequent babies tend to arrive earlier. However, data reveals that the average pregnancy length for first-time births is 38.601 weeks, compared to 38.523 weeks for subsequent births—a difference of about 13 hours. This negligible variation indicates that the perception of first babies arriving significantly later is largely a myth. In reality, their timelines are nearly identical to those of later pregnancies.

## Exercise 2 - 4

```
In [21]: import matplotlib.pyplot as plt
         from math import sqrt

In [22]: # get first born data from pregnancy dataset
         weight_firstborns = preg[preg.birthord == 1]

         # get non first born data from pregnancy data
         weight_others = preg[preg.birthord != 1]
```

```

In [23]: # create histogram of first born total weight
firstborns_hist = weight_firstborns.hist(column = 'totalwgt_lb', bins = 11)

firstborns_hist = firstborns_hist[0]
for x in firstborns_hist:
    # Draw horizontal axis lines
    vals = x.get_yticks()
    for tick in vals:
        x.axhline(y=tick, linestyle='dashed', alpha=0.4, color='#eeeeee',

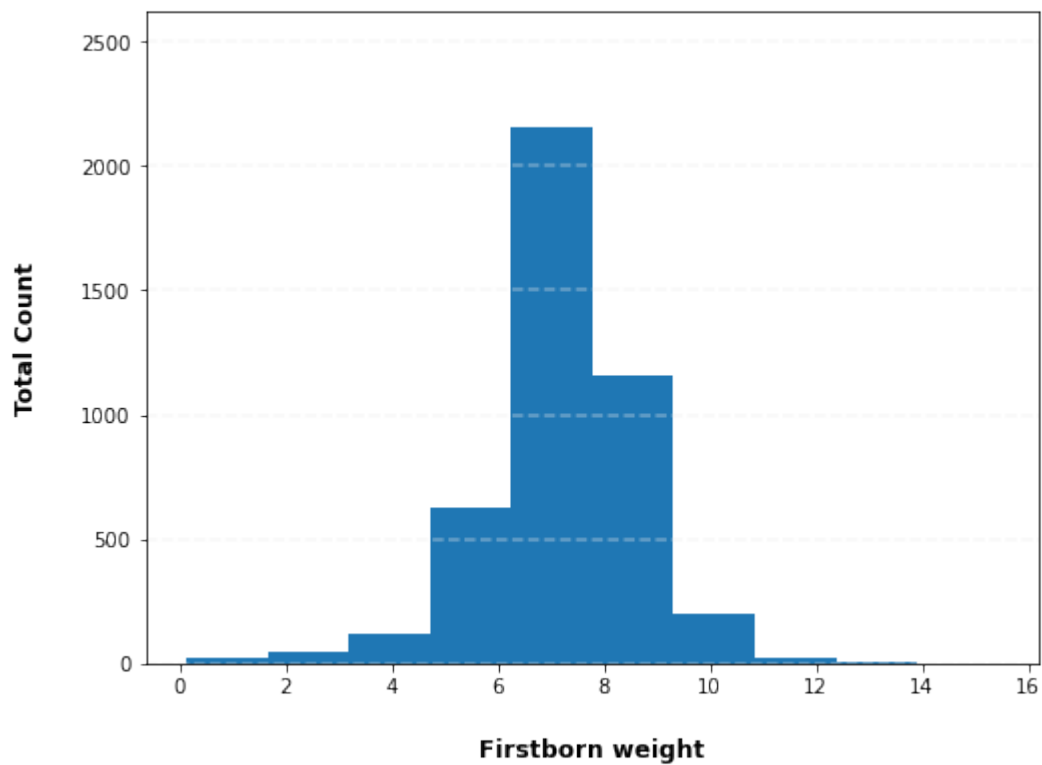
    # Remove title
    x.set_title("")

    # Set X-axis label
    x.set_xlabel("Firstborn weight", labelpad=20, weight='bold', size=12)

    # Set Y-axis label
    x.set_ylabel("Total Count", labelpad=20, weight='bold', size=12)

plt.show()

```





```

In [24]: # create histogram of non-first born total weight
others_hist = weight_others.hist(column = 'totalwgt_lb', bins = 10, figsi:

others_hist = others_hist[0]
for x in others_hist:
    # Draw horizontal axis lines
    vals = x.get_yticks()
    for tick in vals:
        x.axhline(y=tick, linestyle='dashed', alpha=0.4, color='#eeeeee',

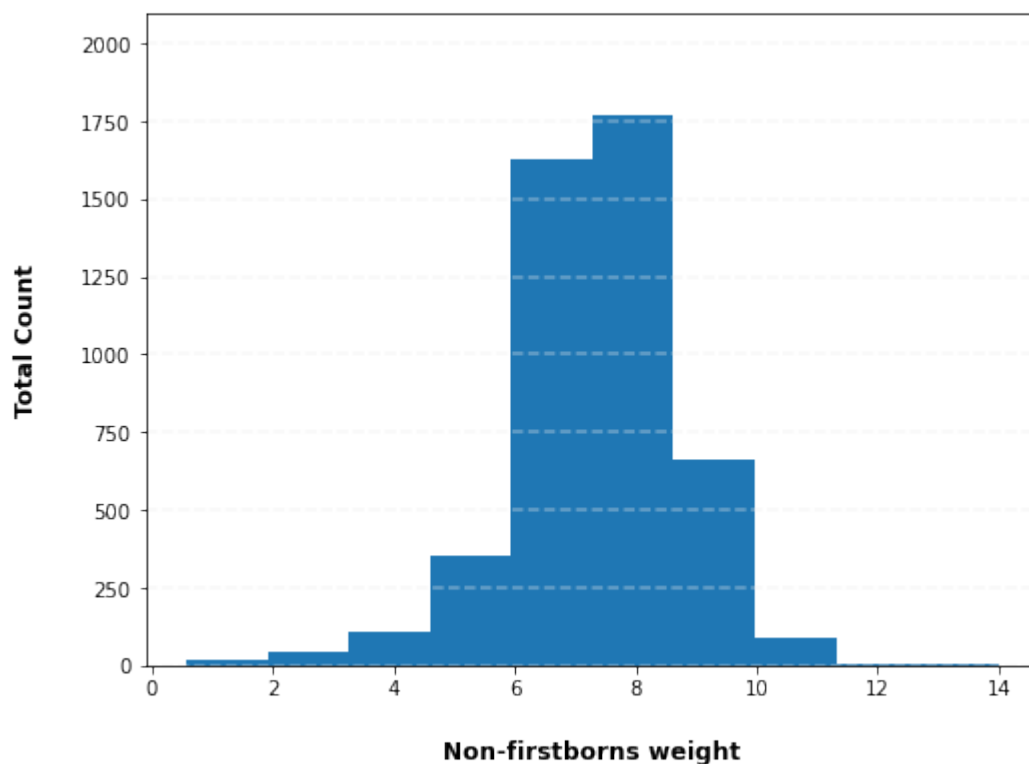
    # Remove title
    x.set_title("")

    # Set X-axis label
    x.set_xlabel("Non-firstborns weight", labelpad=20, weight='bold', size

    # Set Y-axis label
    x.set_ylabel("Total Count", labelpad=20, weight='bold', size=12)

plt.show()

```



```

In [25]: # calculate mean of first born total weight
weight_firstborns_mean = weight_firstborns['totalwgt_lb'].mean()
weight_firstborns_mean

```

Out[25]: 7.201094430437772

```

In [26]: # calculate mean of non-first born total weight
weight_others_mean = weight_others['totalwgt_lb'].mean()
weight_others_mean

```

Out[26]: 7.325855614973262

```

In [29]: def CohensDEffectSize (group1, group2):
    """
    Calculate Cohen's D effect size

    args
        group1 (df): first dataframe to compare
        group2 (df): second dataframe to compare

    returns:
        cohen_d (float): Cohen's D
    """
    meanDiff = group1.mean() - group2.mean()

    var_1 = group1.var()
    var_2 = group2.var()
    len_1, len_2 = len(group1), len(group2)

    pooled_var = (len_1 * var_1 + len_2 * var_2) / (len_1 + len_2)
    cohen_d = meanDiff / sqrt(pooled_var)

    return cohen_d

```

```

In [30]: CohensDEffectSize(weight_firstborns['totalwgt_lb'], weight_others['totalwgt_lb'])

```

```

Out[30]: -0.08893641177719079

```

Total weight for first borns is too small to say that there is a difference at all. With pregnancy length 0.078 which is small also.