

Assignment 8.2

Exercises 9 - 1 and 10 - 1

<http://thinkstats2.com>

Copyright 2016 Allen B. Downey

MIT License: <https://opensource.org/licenses/MIT>

```
In [1]: # Imports
import numpy as np
import matplotlib.pyplot as plt
import thinkstats2
import first

In [2]: from IPython.core.display import HTML
table_css = 'table {align:left;display:block} '
HTML('<style>{}</style>'.format(table_css))

Out[2]:
```

Exercise 9 - 1

As sample size increase, the power of a hypotheses test increase, which means it is more likely to be positive if the effect is real. Conversely, as sample size decreases, the test is less likely to be positive even if the effect is real.

To investigate this behavior, run the tests in this chapter with different subsets of the NSFG data. You can use `thinkstats2.SampleRows` to select a random subset of the rows in a `DataFrame`.

What happens to the p -values of these tests as the sample size decreases? What is the smallest sample size that yields a positive test?

```
In [3]: # Read NSFG data
live_df, firsts_df, others_df = first.MakeFrames()
```

```

In [4]: def PregLengths(sample, iters):
    '''
    Pregnancy lengths: Mean Difference, P-Value

    args:
        sample (df)
        iters (int)

    returns:
        none
    '''
    # Get first borns and others from the sample
    firsts = sample[sample.birthord == 1]
    others = sample[sample.birthord != 1]

    # Get pregnancy lengths from each df and drop any empty rows
    first_prglengths = firsts.prglength.dropna().values
    others_prglengths = others.prglength.dropna().values

    # get the mean difference
    first_mean = first_prglengths.mean()
    others_mean = others_prglengths.mean()
    mean_diff = first_mean - others_mean

    print('\nMean Difference in Preg Length: {} weeks'.format(mean_diff))

    # see probability of this randomly happening
    count = 0

    # loop from 1 to number of trials
    for _ in range(iters):
        # get length of pregnancy lengths
        first_len = len(first_prglengths)
        others_len = len(others_prglengths)

        # get data from sample pregnancy lengths
        all_prglengths = sample.prglength.values
        # shuffle the contents of all pregnancy lengths
        np.random.shuffle(all_prglengths)

        # get pregnancy length from beginning to length of first borns
        new_first = all_prglengths[:first_len]
        # get pregnancy length from length of first borns to end
        new_others = all_prglengths[first_len:]
        # calculate new mean difference
        new_mean_diff = new_first.mean() - new_others.mean()

        # if new mean difference / mean difference >= 1 add 1 to count
        if new_mean_diff / mean_diff >= 1:
            count += 1

    # calculate p-value
    pValue = count / iters

    print(f'P Value: {pValue}')

def BirthWeights(sample, iters):
    '''
    Birth Weights: Mean Difference, P-Value

```

```

In [5]: # Statistical tests - 100 rows
        rows = 100
        trials = 1000

        sample = thinkstats2.SampleRows(live_df, rows)

        print('Sample Size: 100')

        PregLengths(sample, trials)
        BirthWeights(sample, trials)
        Correlate(sample, trials)

        Sample Size: 100

        Mean Difference in Preg Length: 0.20488195278110766 weeks
        P Value: 0.308

        Mean Difference in Birth Weights: 0.30334793491864787 lbs
        P Value: 0.0

        Correlation Coefficient between Age and Weight: 0.0848341795247848
        P Value: 0.196

In [6]: # Statistical tests - 50 rows
        rows = 50
        trials = 1000

        sample = thinkstats2.SampleRows(live_df, rows)

        print('Sample Size: 50')

        PregLengths(sample, trials)
        BirthWeights(sample, trials)
        Correlate(sample, trials)

        Sample Size: 50

        Mean Difference in Preg Length: 0.4614121510673286 weeks
        P Value: 0.289

        Mean Difference in Birth Weights: 0.5591133004926112 lbs
        P Value: 0.069

        Correlation Coefficient between Age and Weight: -0.004576981667129779
        P Value: 0.474

In [7]: # Statistical tests - 10 rows
        rows = 10
        trials = 1000

        sample = thinkstats2.SampleRows(live_df, rows)

        print('Sample Size: 10')

        PregLengths(sample, trials)
        BirthWeights(sample, trials)
        Correlate(sample, trials)

```

Sample Size: 10

Mean Difference in Preg Length: 1.333333333333357 weeks
P Value: 0.306

Mean Difference in Birth Weights: -0.3020833333333304 lbs
P Value: 0.0

Correlation Coefficient between Age and Weight: -0.2719457369136887
P Value: 0.202

test	Sample Size	p-value
Mean Diff in Preg Length	100	0.008
Mean Diff in Preg Length	50	0.006
Mean Diff in Preg Length	10	0.706
Mean Diff in Birth Weight	100	0.442
Mean Diff in Birth Weight	50	0.0
Mean Diff in Birth Weight	10	0.0
Corr Coef: Age and Weight	100	0.211
Corr Coef: Age and Weight	50	0.1
Corr Coef: Age and Weight	10	0.355

The p-values dont really move in a pattern except birth weight which gets down towards 0.

```

In [8]: # chi squared test
n = 100
sample = thinkstats2.SampleRows(live_df, n)

firsts = sample[sample.birthord == 1]
others = sample[sample.birthord != 1]

chi_sq = ChiSquared(live_df, firsts)

# generate null hypothesis data
iters = 1000
chis = []

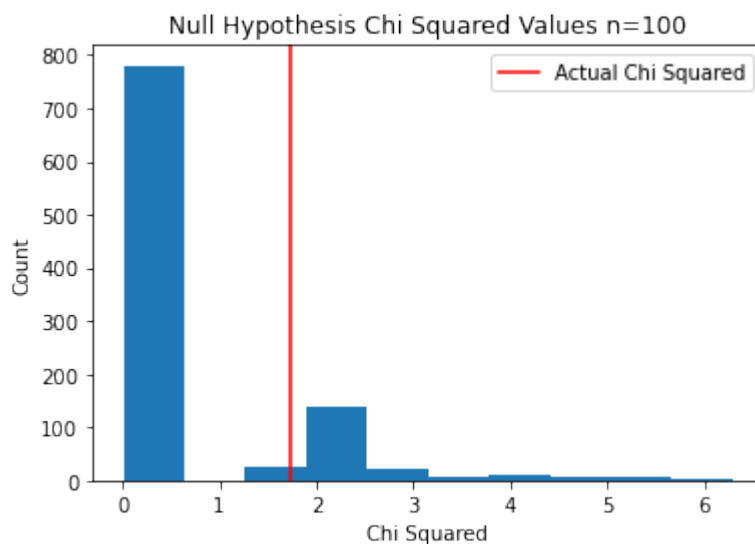
for _ in range(iters):
    samp = thinkstats2.SampleRows(live_df, n)
    chis.append(ChiSquared(live_df, samp))

p_val = len(np.array(chis)[np.where(chis >= chi_sq)]) / len(chis)

plt.hist(chis)
plt.axvline(chi_sq, label='Actual Chi Squared', color='r')
plt.title('Null Hypothesis Chi Squared Values n={}'.format(n))
plt.ylabel('Count')
plt.xlabel('Chi Squared')
plt.legend()
plt.show()

print(f'Max Value Under Null Hypothesis: {max(chis)}')
print(f'Actual Chi Squared: {chi_sq}')
print(f'p-value: {p_val}')

```



```

Max Value Under Null Hypothesis: 6.293512417161693
Actual Chi Squared: 1.7261786610411627
p-value: 0.214

```

```

In [9]: n = 50
        sample = thinkstats2.SampleRows(live_df, n)

        firsts = sample[sample.birthord == 1]
        others = sample[sample.birthord != 1]

        chi_sq = ChiSquared(live_df, firsts)

        # generate null hypothesis data
        iters = 1000
        chis = []


        for _ in range(iters):
            samp = thinkstats2.SampleRows(live_df, n)
            chis.append(ChiSquared(live_df, samp))

        p_val = len(np.array(chis)[np.where(chis >= chi_sq)]) / len(chis)

        plt.hist(chis)
        plt.axvline(chi_sq, label='Actual Chi Squared', color='r')
        plt.title('Null Hypothesis Chi Squared Values n={}'.format(n))
        plt.ylabel('Count')
        plt.xlabel('Chi Squared')
        plt.legend()
        plt.show()

        print(f'Max Value Under Null Hypothesis: {max(chis)}')
        print(f'Actual Chi Squared: {chi_sq}')
        print(f'p-value: {p_val}')

```



```

Max Value Under Null Hypothesis: 10.793459600498872
Actual Chi Squared: 2.6623909883909604
p-value: 0.203

```

```

In [10]: n = 10
         sample = thinkstats2.SampleRows(live_df, n)

         firsts = sample[sample.birthord == 1]
         others = sample[sample.birthord != 1]

         chi_sq = ChiSquared(live_df, firsts)

         # generate null hypothesis data
         iters = 1000
         chis = []

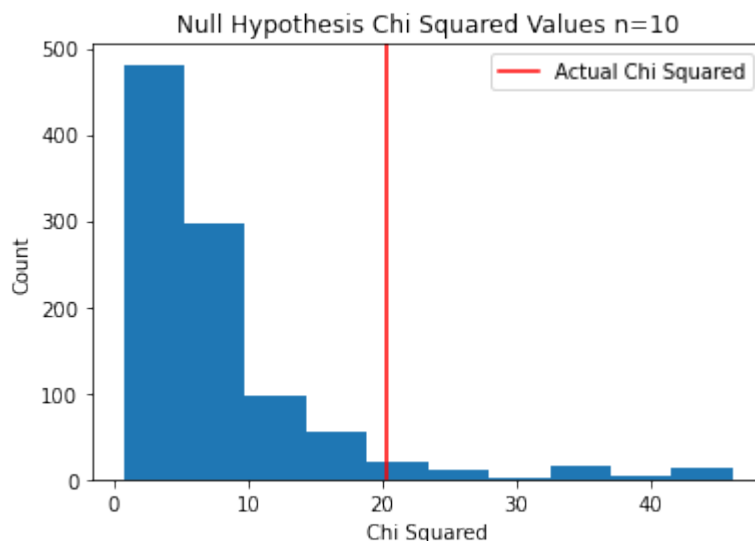
         for _ in range(iters):
             samp = thinkstats2.SampleRows(live_df, n)
             chis.append(ChiSquared(live_df, samp))

         p_val = len(np.array(chis)[np.where(chis >= chi_sq)]) / len(chis)

         plt.hist(chis)
         plt.axvline(chi_sq, label='Actual Chi Squared', color='r')
         plt.title('Null Hypothesis Chi Squared Values n={}'.format(n))
         plt.ylabel('Count')
         plt.xlabel('Chi Squared')
         plt.legend()
         plt.show()

         print(f'Max Value Under Null Hypothesis: {max(chis)}')
         print(f'Actual Chi Squared: {chi_sq}')
         print(f'p-value: {p_val}')

```



```

Max Value Under Null Hypothesis: 46.15692057008778
Actual Chi Squared: 20.35544668462675
p-value: 0.067

```

The p-values for chi-squared are also not following a pattern so it is hard to predict the lowest sample size that yields a positive test.

Exercise 10-1

Using the data from the BRFSS, compute the linear least squares for $\log(\text{weight})$ vs height. How would you best present the estimated parameters for a model like this where one of the variables is transformed? If you were trying to guess someone's weight, how much would it help to know their height?

Like the NSFG, the BRFSS oversamples some groups and provides a sampling weight for each respondent. In the BRFSS data, the variable name for these weights is `totalwt`. Use resampling, without weights, to estimate the mean height of respondents in the BRFSS, the standard error of mean, and a 90% confidence interval. How much does correct weighting affect the estimates?

```
In [11]: import brfss
import scipy
```

```
In [12]: # Read BRFSS data
brfss_df = brfss.ReadBrfss(nrows=None)
brfss_df.head()
```

```
Out[12]:
```

	age	sex	wtyr ago	finalwt	wtkg2	htm3
0	82.0	2	76.363636	185.870345	70.91	157.0
1	65.0	2	72.727273	126.603027	72.73	163.0
2	48.0	2	NaN	181.063210	NaN	165.0
3	61.0	1	73.636364	517.926275	73.64	170.0
4	26.0	1	88.636364	1252.624630	88.64	185.0

```
In [13]: # drop the empty rows from the htm3 and wtkg2 columns
brfss_df = brfss_df.dropna(subset=['htm3', 'wtkg2'])

# move the columns into their own df
heights, weights = brfss_df.htm3, brfss_df.wtkg2
```

```
In [14]: # convert weights to log 10
log_weights = np.log10(weights)
```

```
In [15]: # calculate the linear regression for heights and weights
regress = scipy.stats.linregress(heights, log_weights)
slope, intercept, r_value, p_value, std_err = regress
```

```
print(f"Slope: {slope}")
print(f"Intercept: {intercept}")
print(f"R Value - Pearson: {r_value}")
print(f"P Value: {p_value}")
print(f"Std Error: {std_err}")
```

```
Slope: 0.0052814541694177755
Intercept: 0.9930804163932879
R Value - Pearson: 0.5317282605983431
P Value: 0.0
Std Error: 1.3370549498879916e-05
```



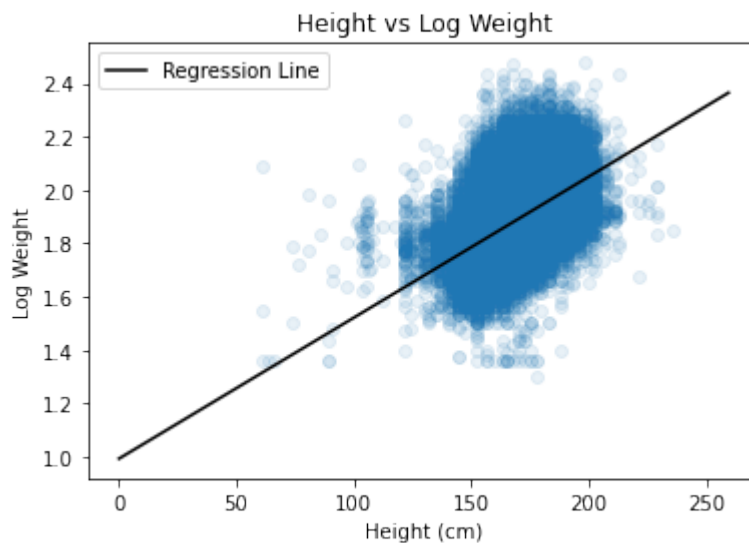
```
In [16]: # Plot regression with log
```

```
plt.scatter(heights, log_weights, alpha = 0.1)

x_fit = np.linspace(0, max(heights)*1.1, 100)
y_fit = slope * x_fit + intercept

plt.plot(x_fit, y_fit, color = 'k', label = 'Regression Line')
plt.title('Height vs Log Weight')
plt.xlabel('Height (cm)')
plt.ylabel('Log Weight')

plt.legend()
plt.show()
```




```
In [17]: # Plot regression without log
```

```
plt.scatter(heights, weights, alpha = 0.1)

x_fit = np.linspace(0, max(heights)*1.1, 100)
y_fit = slope * x_fit + intercept

plt.plot(x_fit, 10**y_fit, color = 'k', label = 'Regression Line')
plt.title('Height vs Weight')
plt.xlabel('Height (cm)')
plt.ylabel('Weight (kg)')

plt.legend()
plt.show()
```



```
In [18]: # Calculate the residuals
residuals = []
```

```
for h, w in zip(heights, log_weights):
    r = w - (h * slope + intercept)
    residuals.append(r)
```

```
In [19]: # height range
height_range = range(140, 190)
```

```
In [20]: # calculate percentiles
```

```
r_25 = []
r_50 = []
r_75 = []
```

```
residuals = np.array(residuals)
```

```
for h in height_range:
    r = residuals[np.where(heights == h)]

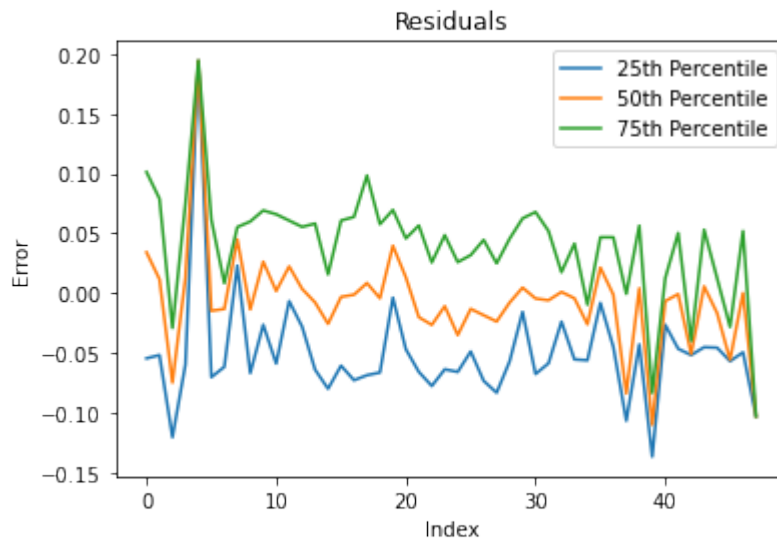
    if r.size != 0:
        r_25.append(np.percentile(residuals[np.where(heights == h)], 25))
        r_50.append(np.percentile(residuals[np.where(heights == h)], 50))
        r_75.append(np.percentile(residuals[np.where(heights == h)], 75))
```

```
In [21]: # Plot percentiles
```

```
plt.plot(r_25, label='25th Percentile')
plt.plot(r_50, label='50th Percentile')
plt.plot(r_75, label='75th Percentile')
```

```
plt.title('Residuals')
plt.xlabel('Index')
plt.ylabel('Error')
```

```
plt.legend()
plt.show()
```



In [22]: # calculate coefficient of determination

```
r2 = thinkstats2.CoeffDetermination(log_weights, residuals)
print(f"r2: {r2}")
```

r2: 0.282734943118939

In [23]: # Confirm that $R^2 = p^2$
 np.isclose(r_value**2, r2)

Out[23]: True

In [24]: # unweighted resampling

```
# resample rows
estimates_unweighted = [thinkstats2.ResampleRows(brfss_df).htm3.mean()
                        for _ in range(100)]
```

```
mean = np.mean(estimates_unweighted)
stderr = np.std(estimates_unweighted)
cdf = thinkstats2.Cdf(estimates_unweighted)
ci = cdf.ConfidenceInterval(90)
```

```
print(f'Mean: {mean}')
print(f'Std Error: {stderr}')
print(f'CI: {ci}')
```

Mean: 168.9567373279573
 Std Error: 0.015794354242453464
 CI: (168.93093787263282, 168.9810879362962)

```

In [25]: # weighted resampling
def ResampleRowsWeighted(df):
    weights = df['finalwt']
    cdf = thinkstats2.Cdf(dict(weights))
    indices = cdf.Sample(len(weights))
    sample = df.loc[indices]
    return sample

# resample rows
estimates_weighted = [ResampleRowsWeighted(brfss_df).htm3.mean()
                      for _ in range(100)]

mean = np.mean(estimates_weighted)
stderr = np.std(estimates_weighted)
cdf = thinkstats2.Cdf(estimates_weighted)
ci = cdf.ConfidenceInterval(90)

print(f'Mean: {mean}')
print(f'Std Error: {stderr}')
print(f'CI: {ci}')

Mean: 170.4965447209928
Std Error: 0.016775810759808673
CI: (170.46839315669274, 170.52121101881608)

```

With weighting the mean difference is only about 1.5 kg with a standard error difference of only a 0.004 which implies that the difference is not due to chance.